

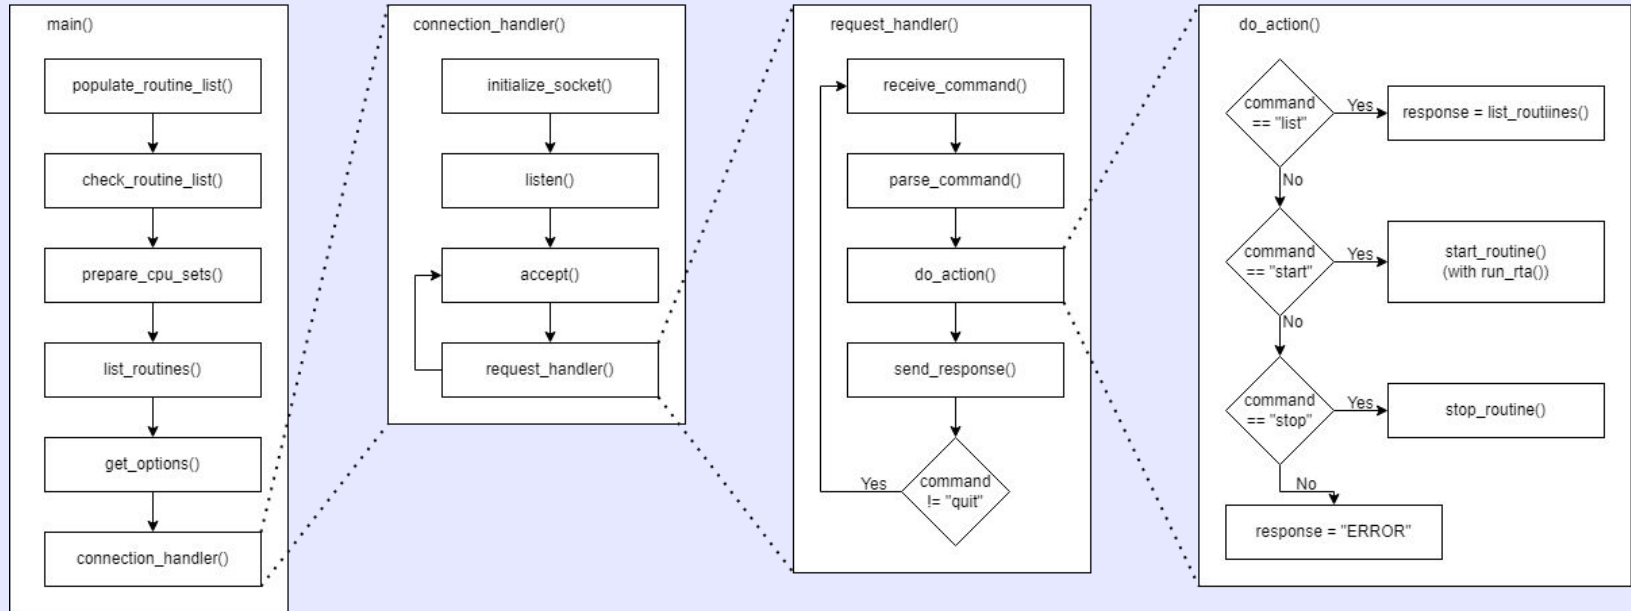


Concurrent and Real-Time Programming

Exercise N°5

Lorenzo Cappellotto
2044728

General Structure (Supervisor)



Used Data Structures

```

/*****
// Static variables and MACRO definitions

#define MAX_LENGTH_ROUTINE_NAME 31
#define MAX_NUMBER_ROUTINES 98
#define MAX_NUMBER_CORES 8
#define DEFAULT_CORE_NUMBER 0
#define BILLION 1000000000.0
#define WTA_NUMBER_RUNS 50 // 1K

static const char short_options[] = "p:htsm";

static struct option long_options[] = {
    {"port",                required_argument, 0, 'p'},
    {"help",                no_argument, 0, 'h'},
    {"time-analysis",       no_argument, 0, 't'},
    {"multicore-scheduling", no_argument, 0, 'm'},
    {0, 0, 0, 0}
};

static int port = 50000;
static int sd, currSd;
static int reuse = 1;

// mode_multicore == 0 means single-core, mode_multicore == 1 means multi-core.
static int mode_multicore = 0;

```

```

static socklen_t accept_sin_length;
static struct sockaddr_in listen_sin, accept_sin;

static int number_routines = 0;

static struct routine {
    char routine_name[MAX_LENGTH_ROUTINE_NAME];
    void *(*routine_pointer)(void*);
    double period;
    double deadline;
    double wcet;
    double utilization;
    int active_running_core[MAX_NUMBER_CORES];
} routines_list[MAX_NUMBER_ROUTINES];

static struct routine_command {
    char *action;
    char *routine_name;
    int core_number;
    int routine_index;
} parsed_command;

static char *core_number_string;

```

Used Data Structures and CPU_SET Preparation

```
routine_listings:
routine: 'fun1', pointer: 39e3f3e9, worst_execution_time: 3.0, period: 8.0,
deadline: 8.0, utilization: 0.375
    active cores: 1 1 0
routine: 'fun2', pointer: 39e3f440, worst_execution_time: 4.0, period: 14.0,
deadline: 14.0, utilization: 0.286
    active cores: 0 1 0
routine: 'fun3', pointer: 39e3f497, worst_execution_time: 5.0, period: 22.0,
deadline: 22.0, utilization: 0.227
    active cores: 0 0 1
routine: 'fun4', pointer: 39e3f4ee, worst_execution_time: 7.0, period: 30.0,
deadline: 30.0, utilization: 0.233
    active cores: 0 0 1
```

core_set		threads	Number of routines					
0	001		0,0	0,1	0,2	0,3	...	0,98
1	010	Number of cores	1,0	1,1	1,2	1,3	...	1,98
2	100		2,0	2,1	2,2	2,3	...	2,98

Core Bit
Number Mask

```
static pthread_t threads[MAX_NUMBER_CORES][MAX_NUMBER_ROUTINES];
static cpu_set_t core_set[MAX_NUMBER_CORES];
static long number_of_cores;
static pthread_attr_t routine_attr;

static int max_priority;
static char *routines_listing;

/*****/

static void prepare_cpu_sets()
{
    number_of_cores = sysconf(_SC_NPROCESSORS_ONLN);
    max_priority = sched_get_priority_max(SCHED_FIFO);
    printf("INFO: The maximum priority for SCHED_FIFO is %d\n", max_priority);

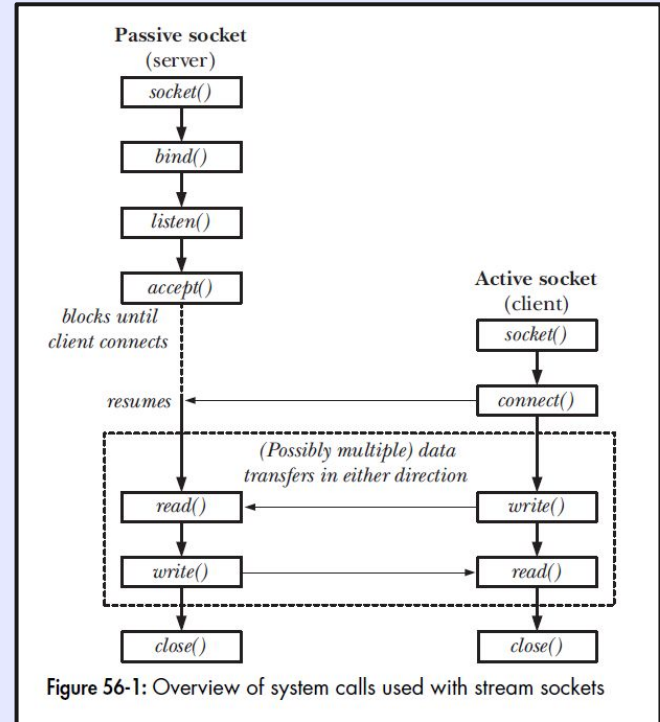
    if (number_of_cores > MAX_NUMBER_CORES)
        perror_exit("ERROR: the real number of cores is greater than the "
            "maximum number\n");

    for (size_t i = 0; i < number_of_cores; i++) {
        CPU_ZERO(&core_set[i]);
        CPU_SET(i, &core_set[i]);
    }
}
```

Client/Server Model

```
static void connection_handler()
{ // Shortened version ...
    if (-1 == (sd = socket(AF_INET, SOCK_STREAM, 0)))
        perror_exit("ERROR: The system could not create the socket\n");
    if (-1 == setsockopt(sd, SOL_SOCKET, SO_REUSEADDR, (const char*)&reuse,
        sizeof(reuse)))
        perror_exit("ERROR: setsockopt(SO_REUSEADDR) failed\n");
    if (-1 == setsockopt(sd, SOL_SOCKET, SO_REUSEPORT, (const char*)&reuse,
        sizeof(reuse)))
        perror_exit("ERROR: setsockopt(SO_REUSEPORT) failed\n");
    memset(&listen_sin, 0, sizeof(listen_sin));
    listen_sin.sin_family = AF_INET;
    listen_sin.sin_addr.s_addr = INADDR_ANY;
    listen_sin.sin_port = htons(port);

    if (-1 == bind(sd, (struct sockaddr *)&listen_sin, sizeof(listen_sin)))
        perror_exit("ERROR: The system could not bind the socket\n");
    if (-1 == listen(sd, 5))
        perror_exit("ERROR: listen() was unsuccessful on the socket.");
    accept_sin_length = sizeof(accept_sin);
    for(;;)
        if (-1 != (currSd = accept(sd, (struct sockaddr *) &accept_sin,
            &accept_sin_length)))
            request_handler(currSd);
    close(sd);
}
```



Client/Server Model

```
static void request_handler(int currSd)
{ // Shortened version.
    unsigned int network_string_length;
    int pcr, string_length, exit_status = 0;
    char *command, *response;

    for(;;) {
        if (-1 == receive(currSd, (char *)&network_string_length,
            sizeof(network_string_length)))
            break;
        string_length = ntohs(network_string_length);
        command = malloc(string_length+1);
        if (-1 == receive(currSd, command, string_length))
            break;
        command[string_length] = 0;
        printf("INFO: The command received is '%s'\n", command);
        // If command is "quit" or "list" do ...
        if (0 == strcmp(command, "quit")) { // ...
        } else if (0 == strcmp(command, "list")) { // ...
        } else {
            // The command is of type: "start/stop routine_name core_number"
            if (0 == (pcr = parse_command(&parsed_command, command,
                string_length))) {
                if (0 == strncmp(parsed_command.action, "stop", 4)) {
                    if (-1 == stop_routine())
                        response = strdup("The routine could not be stopped");
```

```
                    else
                        response = strdup("The routine has stopped correctly");
                } else {
                    if (-1 == start_routine())
                        response = strdup("The routine could not be started");
                    else
                        response = strdup("The routine has started correctly");
                }
            } else {
                response = strdup("The command given is malformed");
            }
        }
        printf("INFO: The response is '%s'\n", response);
        string_length = strlen(response);
        network_string_length = htons(string_length);
        if (-1 == send(currSd, &network_string_length,
            sizeof(network_string_length), 0))
            break;
        if (-1 == send(currSd, response, string_length, 0))
            break;
        free(command); free(response);
        if (exit_status)
            break;
    }
    close(currSd);
}
```

```

network_string_length = htonl(network_string_length);
if (-1 == send(sd, &network_string_length,
               sizeof(network_string_length), 0))
    perror_exit("ERROR: Impossible to send the length");
if (-1 == send(sd, command, string_length, 0))
    perror_exit("ERROR: Impossible to send the command");

if (-1 == receive(sd, (char *)&network_string_length,
                  sizeof(network_string_length)))
    perror_exit("ERROR: Impossible to receive the length");
string_length = ntohl(network_string_length);
response = malloc(string_length + 1);
if (-1 == receive(sd, response, string_length))
    perror_exit("ERROR: Impossible to receive the response");
response[string_length] = 0;
printf("INFO: The response is '%s'\n", response);
free(response);

if(0 == strcmp(command, "quit"))
    break;
}

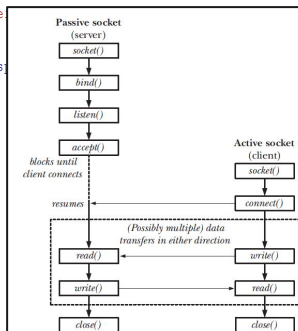
close(sd);
printf("INFO: Connection Terminated\n");
return 0;
}

```

```

graph TD
    A[socket()] --> B[bind()]
    B --> C[listen()]
    C --> D[accept()]
    D -- "blocks until client connects" --> E[read()]
    E -- "(Possibly transfers in)" --> F[write()]
    F --> G[close()]
    style D stroke-width:2px
    subgraph " "
        E
        F
    end

```



Deadline Monotonic Scheduling

Index	0	1	2	3	...	98
RT_prio	99-0	99-1	97	96	...	1
PR (-1-RT_prio)	-100	-99	-98	-97	...	-2

```
static int stop_routine()
{ // Shortened version.
    if (0 == routines_list[parsed_command.routine_index].
        active_running_core[parsed_command.core_number])
        return -1;
    pthread_cancel(
        threads[parsed_command.core_number][parsed_command.routine_index]);
    pthread_join(
        threads[parsed_command.core_number][parsed_command.routine_index],
        NULL);

    routines_list[parsed_command.routine_index].
        active_running_core[parsed_command.core_number] = 0;
    return 0;
}
```

```
static int set_realtime_attribute(pthread_attr_t *attr, int policy,
    int priority, cpu_set_t *cpuset) { // Shortened version.
    struct sched_param param;
    int status;
    pthread_attr_init(attr);
    status = pthread_attr_getschedparam(attr, &param);
    if(status)
        return status;
    status = pthread_attr_setinheritsched(attr, PTHREAD_EXPLICIT_SCHED);
    if(status)
        return status;
    status = pthread_attr_setschedpolicy(attr, policy);
    if(status)
        return status;
    param.sched_priority = priority;
    status = pthread_attr_setschedparam(attr, &param);
    if(status)
        return status;
    if(cpuset != NULL) {
        status = pthread_attr_setaffinity_np(attr, sizeof(cpu_set_t), cpuset);
        if(status)
            return status;
    }
    return status;
}
```


Deadline Monotonic Scheduling

```
static int start_routine()
{ // Shortened version.
    int schedulable;
    if (1 == routines_list[parsed_command.routine_index].
        active_running_core[parsed_command.core_number])
        return -1;
    if (-1 == (schedulable =
        run_rta(parsed_command.routine_index, parsed_command.core_number)))
        return -1;

    pthread_attr_t routine_attr;
    set_realtime_attribute(
        &routine_attr,
        SCHED_FIFO,
        max_priority - parsed_command.routine_index,
        &core_set[parsed_command.core_number]);
    pthread_create(
        &threads[parsed_command.core_number][parsed_command.routine_index],
        &routine_attr,
        routine_wrapper,
        &routines_list[parsed_command.routine_index]);

    routines_list[parsed_command.routine_index].
        active_running_core[parsed_command.core_number] = 1;
    return 0;
}
```

```
$ top -H -p $(ps -a | grep 'supervisor' | awk -F' ' '{print $1}')
```

Fields Management for window **1:Def**, whose current sort field is **%CPU**

Navigate with Up/Dn, Right selects for move then <Enter> or Left commits, 'd' or <Space> toggles display, 's' sets sort. Use 'q' or <Esc> to end!

* PID	= Process Id	GROUP	= Group Name	OOMS	= OOMEM Score c
* USER	= Effective Use	PGRP	= Process Group	ENVIRON	= Environment v
* PR	= Priority	TTY	= Controlling T	vmj	= Major Faults
* NI	= Nice Value	TPGID	= Tty Process G	vmn	= Minor Faults
* VIRT	= Virtual Image	SID	= Session Id	USED	= Res+Swap Size
* RES	= Resident Size	nTH	= Number of Thr	nsIPC	= IPC namespace
* SHR	= Shared Memory	TIME	= CPU Time	nsMNT	= MNT namespace
* S	= Process Statu	SWAP	= Swapped Size	nsNET	= NET namespace
* P	= Last Used Cpu	CODE	= Code Size (Ki	nsPID	= PID namespace
* %CPU	= CPU Usage	DATA	= Data+Stack (K	nsUSER	= USER namespac
* %MEM	= Memory Usage	nMaj	= Major Page Fa	nsUTS	= UTS namespace
* TIME+	= CPU Time, hun	nMin	= Minor Page Fa	LXC	= LXC container
* COMMAND	= Command Name/	nDRT	= Dirty Pages C	RSan	= RES Anonymous
PPID	= Parent Proces	WCHAN	= Sleeping in F	RSfd	= RES File-base
UID	= Effective Use	Flags	= Task Flags <s	RSlk	= RES Locked (K
RUID	= Real User Id	CGROUPS	= Control Group	RSsh	= RES Shared (K
RUSER	= Real User Nam	SUPGIDS	= Supp Groups I	CGNAME	= Control Group
SUID	= Saved User Id	SUPGRPS	= Supp Groups N	NU	= Last Used NUM
SUSER	= Saved User Na	TGID	= Thread Group		
GID	= Group Id	OOMa	= OOMEM Adjustm		

Response Time Analysis

```
static int run_rta(int routine_index, int core_number)
{ // Shortened version.
    double total_utilization = 0.0, interference_time[number_routines];
    double response_time = 0.0, current_response_time = 0.0;

    // Add temporarily the routine to the chosen core in order
    routines_list[routine_index].active_running_core[core_number] = 1;

    // Compute the total utilization
    // ...

    for (int i = 0; i < number_routines; i++) {
        // Skip the iteration for the routines not active on the core.
        if (0 == routines_list[i].active_running_core[core_number])
            continue;

        response_time = -1.0;
        current_response_time = routines_list[i].wcet;
        printf("routine %d, execution_time: %f\n", i, current_response_time);

        while (current_response_time != response_time) {
            response_time = current_response_time;
            current_response_time = routines_list[i].wcet;
        }
    }
}
```

```
for (int j = 0; j < i; j++) {
    if (0 == routines_list[j].active_running_core[core_number])
        continue;

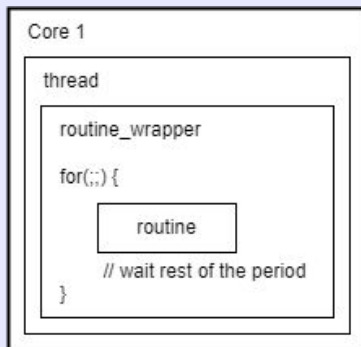
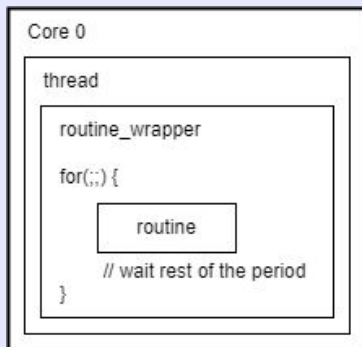
    current_response_time +=
        ceil(response_time / routines_list[j].period) *
        routines_list[j].wcet;
}

if(response_time > routines_list[i].deadline) {
    routines_list[routine_index].active_running_core[core_number] = 0;
    printf("INFO: The set of tasks on core %d is not schedulable using "
        "RM with utilization %.3f\n",
        core_number, total_utilization);
    return -1;
}

routines_list[routine_index].active_running_core[core_number] = 0;
printf("INFO: The set of tasks on core %d is schedulable using RM with "
    "utilization %.3f\n", core_number, total_utilization);
return 0;
}
```

Routine Wrapping and Period Waiting

```
pthread_create(  
    &threads[parsed_command.core_number][parsed_command.routine_index],  
    &routine_attr,  
    routine_wrapper,  
    &routines_list[parsed_command.routine_index]);
```



```
void *fun1(void *args)  
{  
    int x = 1;  
    for(int i = 0; i < 500000000; i++)  
        x *= 1.00001;  
    printf("Fun1\n");  
}  
  
void* routine_wrapper(void *args)  
{  
    struct routine *current_arguments = (struct routine*) args;  
  
    struct timespec remain;  
    int s;  
  
    // An infinite loop, run the routine and wait the rest of the period.  
    for(;;) {  
        clock_gettime(CLOCK_REALTIME, &remain);  
        remain.tv_sec += current_arguments->period;  
        current_arguments->routine_pointer(NULL);  
        s = clock_nanosleep(CLOCK_REALTIME, TIMER_ABSTIME, &remain, NULL);  
    }  
}
```

Multicore Execution

```
// mode_multicore == 0 means single-core, mode_multicore == 1 means multi-core.
static int mode_multicore = 0;

/*****
// Main function

int main (int argc, char **argv)
{ // Shortened version ...
    // ...
    int index_options, c;
    for (;;) {
        c = getopt_long(argc, argv, short_options, long_options,
                        &index_options);
        if (-1 == c)
            break;
        switch (c) {
            // ...
            case 'm':
                mode_multicore = 1;
                printf("INFO: Multicore option selected\n");
                break;
            // ...
        }
    }
}
```

```
static int parse_command(struct routine_command *parsed_command, char *command,
                        int command_length)
{ // Shortened version.

    // Try to parse the action that needs to be done.
    // And check if the action requested actually exists.
    // ...

    // Try to parse the routine name.
    // And find the correct routine associated to the request (if there is one).
    // ...

    // Try to obtain the integer core_number.
    parsed_command->core_number = DEFAULT_CORE_NUMBER;
    core_number_string = strtok(NULL, " ");
    if (mode_multicore == 1 && core_number_string != NULL)
        parsed_command->core_number = atoi(core_number_string);

    // Check the processor number is valid.
    // ...

    return 0;
}
```

Estimating Routines' Worst Case Execution Time

```

/*****
// Main function

int main (int argc, char **argv)
{ // Shortened version ...
    // ...
    int index_options, c;
    for (;;) {
        c = getopt_long(argc, argv, short_options, long_options,
                        &index_options);
        if (-1 == c)
            break;
        switch (c) {
            // ...
            case 't':
                routines_list_wcet_analysis();

                printf("INFO: The updated complete list of routines is:\n");
                list_routines();
                printf("%s", routines_listing);
                break;
            // ...

```

```

static double routine_wcet_analysis(struct routine *myroutine)
{ // Shortened version ...
    double temp, result = 0.0;
    struct timespec start, end;
    for (int i = 0; i < WTA_NUMBER_RUNS; i++) {
        clock_gettime(CLOCK_REALTIME, &start);
        myroutine->routine_pointer(NULL);
        clock_gettime(CLOCK_REALTIME, &end);
        temp = (end.tv_sec - start.tv_sec)
              + (end.tv_nsec - start.tv_nsec) / BILLION;
        if (temp > result)
            result = temp;
    }
    return result;
}

static void routines_list_wcet_analysis()
{ // Shortened version ...
    for (int i = 0; i < number_routines; i++)
        routines_list[i].wcet = routine_wcet_analysis(&routines_list[i]);
    for (int i = 0; i < number_routines; i++) {
        routines_list[i].wcet *= 1.5;
        routines_list[i].utilization =
            routines_list[i].wcet / routines_list[i].period;
    }
}

```



The end

Thank you for your attention!